

## EXP NO: 6: Implement Gradient Descent

Aim:

Implement gradient descent and back-propagation for a deep neural network from scratch (numpy only), train it on a non-linear dataset (KOR), and demonstrate convergence with loss /accuracy.

Objectives:

1. Build a multi-layer perceptron with ReLU hidden layers and softmax input-output.
2. Derive and code forward pass, loss, and backward pass. → analytical gradient. ↗ cross entropy
3. Train with vanilla gradient descent; show loss decreasing and accuracy improving.
4. Reflect on effects of learning rate, initialization, and depth.

Algorithm

1. Initialize parameters:

→ Randomly assign small values to weights  $w$  and set biases  $b$  to zero (or small values).

2. Forward pass:

- Input data goes through each layer:  
→ Multiply input by weight, add

$$\text{bias} \rightarrow z = \text{aprev} \cdot w + b$$

→ Apply activation layer function (like ReLU or sigmoid) to get a

- At last layer, use softmax to get probabilities of each class.

### 3. Compute loss:

- Compare predicted output  $\hat{y}$  with true output  $y$  using a loss function

### 4. Backward pass:

- Calculate error at output layer:  $\delta = \hat{y} - y$
- Multiply error with weight matrix
- Apply derivative of activation function
- Find gradients of weights and biases.

### 5. Update parameters:

- Adjust weights and biases using learning rate  $\eta$ :

$$\rightarrow w = w - \eta \cdot dw$$

$$\rightarrow b = b - \eta \cdot db$$

### 6. Repeat:

- Do step 2 → 5 for many epochs
- Each time, loss decreases and predictions get better.

## 7 Stop Training:

- when loss becomes very small or accuracy is high enough.

## Pseudo Code:

initialize  $w[1], b[1]$  with the init for

$l = 1 \dots L$

for epoch in  $1 \dots E$ :

$a[0] = x$

for  $l$  in  $1 \dots L-1$ :

$$z[l] = a[l-1] @ w[l] + b[l]$$

$$a[l] = \text{ReLU}(z[l])$$

$$z[L] = a[L-1] @ w[L] + b[L]$$

$$y_{\text{hat}} = \text{softmax}(z[L])$$

$$\text{loss} = \text{cross-entropy}(y_{\text{hat}}, y)$$

$$dz[L] = (y_{\text{hat}} - y) / N$$

$$dw[L] = a[L-1]^T @ dz[L]$$

$$db[L] = \text{sum\_rows}(dz[L])$$

$$da = dz[L] @ w[L]^T$$

for  $l$  in  $L-1$  down to  $1$ :

$$dz[l] = da * \text{ReLU}'(z[l])$$

$$dw[l] = a[l-1]^T @ dz[l]$$

$$db[l] = \text{sum\_rows}(dz[l])$$

$$\text{if } l > 1: da = dz[l] @ w[l]^T$$

for  $l$  in  $1 \dots L$ :

$$w[l] = lr * dw[l]$$

$$b[l] \leftarrow lr * db[l]$$

### Observation:

- Convergence: with He initialization + ReLU, vanilla gradient descent learns XOR reliably; loss smoothly decrease to  $\approx 10^{-4}$  and accuracy reaches 100%.
- Learning rate: if  $lr$  is too small, convergence is slow; too large leads to oscillation/divergence. For this setup, 0.05-0.2 works well.
- Depth matters: A single linear layer can't solve XOR; adding at least one non-linear hidden layer is essential. Deeper networks can fit more complex patterns but need good initialization and possibly regularization.
- stability: softmax with log-sum-exp stabilization and one-hot labels keeps training numerically stable.

### Conclusion:-

Gradient Descent with Backpropagation effectively trains a deep neural network. A non-linear hidden layer is essential for solving problems like XOR, and with proper initialization and learning rate, the model converges to accurate results.

[1]

```
▶ import numpy as np
    import matplotlib.pyplot as plt
    np.random.seed(42)
    X_base = np.array([[0,0],[0,1],[1,0],[1,1]], dtype=np.float64)
    y_base = np.array([0,1,1,0])
    repeats = 500
    X = np.tile(X_base, (repeats, 1))
    y = np.tile(y_base, repeats)
    num_classes = 2
    Y = np.eye(num_classes)[y]
    layers = [2, 8, 8, 2]
    lr = 0.1
    epochs = 5000
    print_every = 500
    def he_init(fan_in, fan_out):
        return np.random.randn(fan_in, fan_out) * np.sqrt(2.0 / fan_in)
    def relu(z):
        return np.maximum(0, z)
    def drelu(z):
        return (z > 0).astype(z.dtype)
    def softmax(z):
        z = z - np.max(z, axis=1, keepdims=True)
        exp = np.exp(z)
        return exp / np.sum(exp, axis=1, keepdims=True)
    def cross_entropy(yhat, Y):
        eps = 1e-9
        return -np.mean(Y * np.log(yhat + eps), axis=1)
    def accuracy(yhat, y):
        return np.mean(np.argmax(yhat, axis=1) == y)
    params = {}
    for i in range(len(layers) - 1):
        params[f"W{i+1}"] = he_init(layers[i], layers[i+1])
        params[f"b{i+1}"] = np.zeros((1, layers[i+1]))
    def forward(X, params):
        cache = {}
        a = X
```



Commands + Code + Text ▶ Run all ▾



[1]

```
cache["a0"] = X
L = len(layers) - 1
for i in range(1, L):
    z = a @ params[f"W{i}"] + params[f"b{i}"]
    a = relu(z)
    cache[f"z{i}"] = z
    cache[f"a{i}"] = a
zL = a @ params[f"W{L}"] + params[f"b{L}"]
yhat = softmax(zL)
cache[f"z{L}"] = zL
cache["yhat"] = yhat
return yhat, cache

def backward(Y, cache, params):
    grads = {}
    L = len(layers) - 1
    yhat = cache["yhat"]
    N = Y.shape[0]
    dz = (yhat - Y) / N
    a_prev = cache[f"a{L-1}"]
    grads[f"dw{L}"] = a_prev.T @ dz
    grads[f"db{L}"] = np.sum(dz, axis=0, keepdims=True)
    da_prev = dz @ params[f"W{L}"].T
    for i in range(L-1, 0, -1):
        z = cache[f"z{i}"]
        dz = da_prev * drelu(z)
        a_prev = cache["a0"] if i == 1 else cache[f"a{i-1}"]
        grads[f"dw{i}"] = a_prev.T @ dz
        grads[f"db{i}"] = np.sum(dz, axis=0, keepdims=True)
        if i > 1:
            da_prev = dz @ params[f"W{i}"].T
    return grads

losses = []
accuracies = []
epochs_list = []
for epoch in range(1, epochs + 1):
    yhat, cache = forward(X, params)
    Ls = cross_entropy(yhat, Y)
    grads = backward(Y, cache, params)
```

```
[1] da_prev = dz @ params[f'W{1}'].T  
    return grads  
losses = []  
accuracies = []  
epochs_list = []  
for epoch in range(1, epochs + 1):  
    yhat, cache = forward(X, params)  
    Ls = cross_entropy(yhat, Y)  
    grads = backward(Y, cache, params)  
    for k in params:  
        params[k] -= lr * grads["d" + k]  
    if epoch % print_every == 0:  
        acc = accuracy(yhat, y)  
        losses.append(Ls)  
        accuracies.append(acc)  
        epochs_list.append(epoch)  
        print(f"Epoch {epoch:4d} | loss={Ls:.6f} | acc={acc:.4f}")  
yhat, _ = forward(X, params)  
print("\nFinal accuracy:", accuracy(yhat, y))  
print("Sample predictions (first 8):", np.argmax(yhat[:8], axis=1))  
print("Sample truths      (first 8):", y[:8])  
plt.figure(figsize=(12,5))  
plt.subplot(1,2,1)  
plt.plot(epochs_list, losses, marker='o')  
plt.title('Training Loss over Epochs')  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.grid(True)  
plt.subplot(1,2,2)  
plt.plot(epochs_list, accuracies, marker='o', color='green')  
plt.title('Training Accuracy over Epochs')  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.ylim([0, 1.05])  
plt.grid(True)  
plt.show()
```

[1]  
11s

```
Epoch 500 | loss=0.004159 | acc=1.0000
Epoch 1000 | loss=0.001368 | acc=1.0000
Epoch 1500 | loss=0.000765 | acc=1.0000
Epoch 2000 | loss=0.000516 | acc=1.0000
Epoch 2500 | loss=0.000384 | acc=1.0000
Epoch 3000 | loss=0.000303 | acc=1.0000
Epoch 3500 | loss=0.000248 | acc=1.0000
Epoch 4000 | loss=0.000209 | acc=1.0000
Epoch 4500 | loss=0.000180 | acc=1.0000
Epoch 5000 | loss=0.000158 | acc=1.0000
```

Final accuracy: 1.0

Sample predictions (first 8): [0 1 1 0 0 1 1 0]

Sample truths (first 8): [0 1 1 0 0 1 1 0]

