

EXP NO: 11: Experiments using Variational

Aim:-

To implement a Variational Autoencoder (VAE) on the MNIST dataset and generate new handwritten digits from the learned latent space

Objectives

1. To understand probabilistic representation learning using VAEs
2. To generate new data samples from a learned latent distribution
3. To explore the concept of mean, variance, and reparameterization trick in generative models.
4. To visualize the latent space and generated images.

Algorithm:-

1. Load and normalize MNIST dataset.
2. Define encoder network to output mean (μ) and log variance ($\log \sigma^2$) for latent variables.
3. Apply reparameterization trick

$$z = \mu + \sigma * \epsilon \text{ where } \epsilon \sim \mathcal{N}(0, 1)$$
4. Define decoder network to reconstruct image from latent vector z .
5. Compute VAE loss

$$\text{Loss} = \text{Reconstruction-Loss} + k\text{L-Divergence}$$

where KL divergence regularizes latent space towards a normal distribution

6. Train the VAE model
7. Visualize reconstructed and generated samples.

Pseudo Code:-

Load MNIST dataset

Normalize to $[0, 1]$

Define encoder:

Dense(256, relu)

Dense(latent-mean)

Dense(latent-log-var)

Reparameterization:

$$z = \text{mean} + \exp(\text{log-var}) * \epsilon$$

Define decoder:

Dense(256, relu)

Dense(784, sigmoid)

Compute VAE loss:

reconstruction loss + KL-divergence

Train model

visualize reconstructed and generated images.

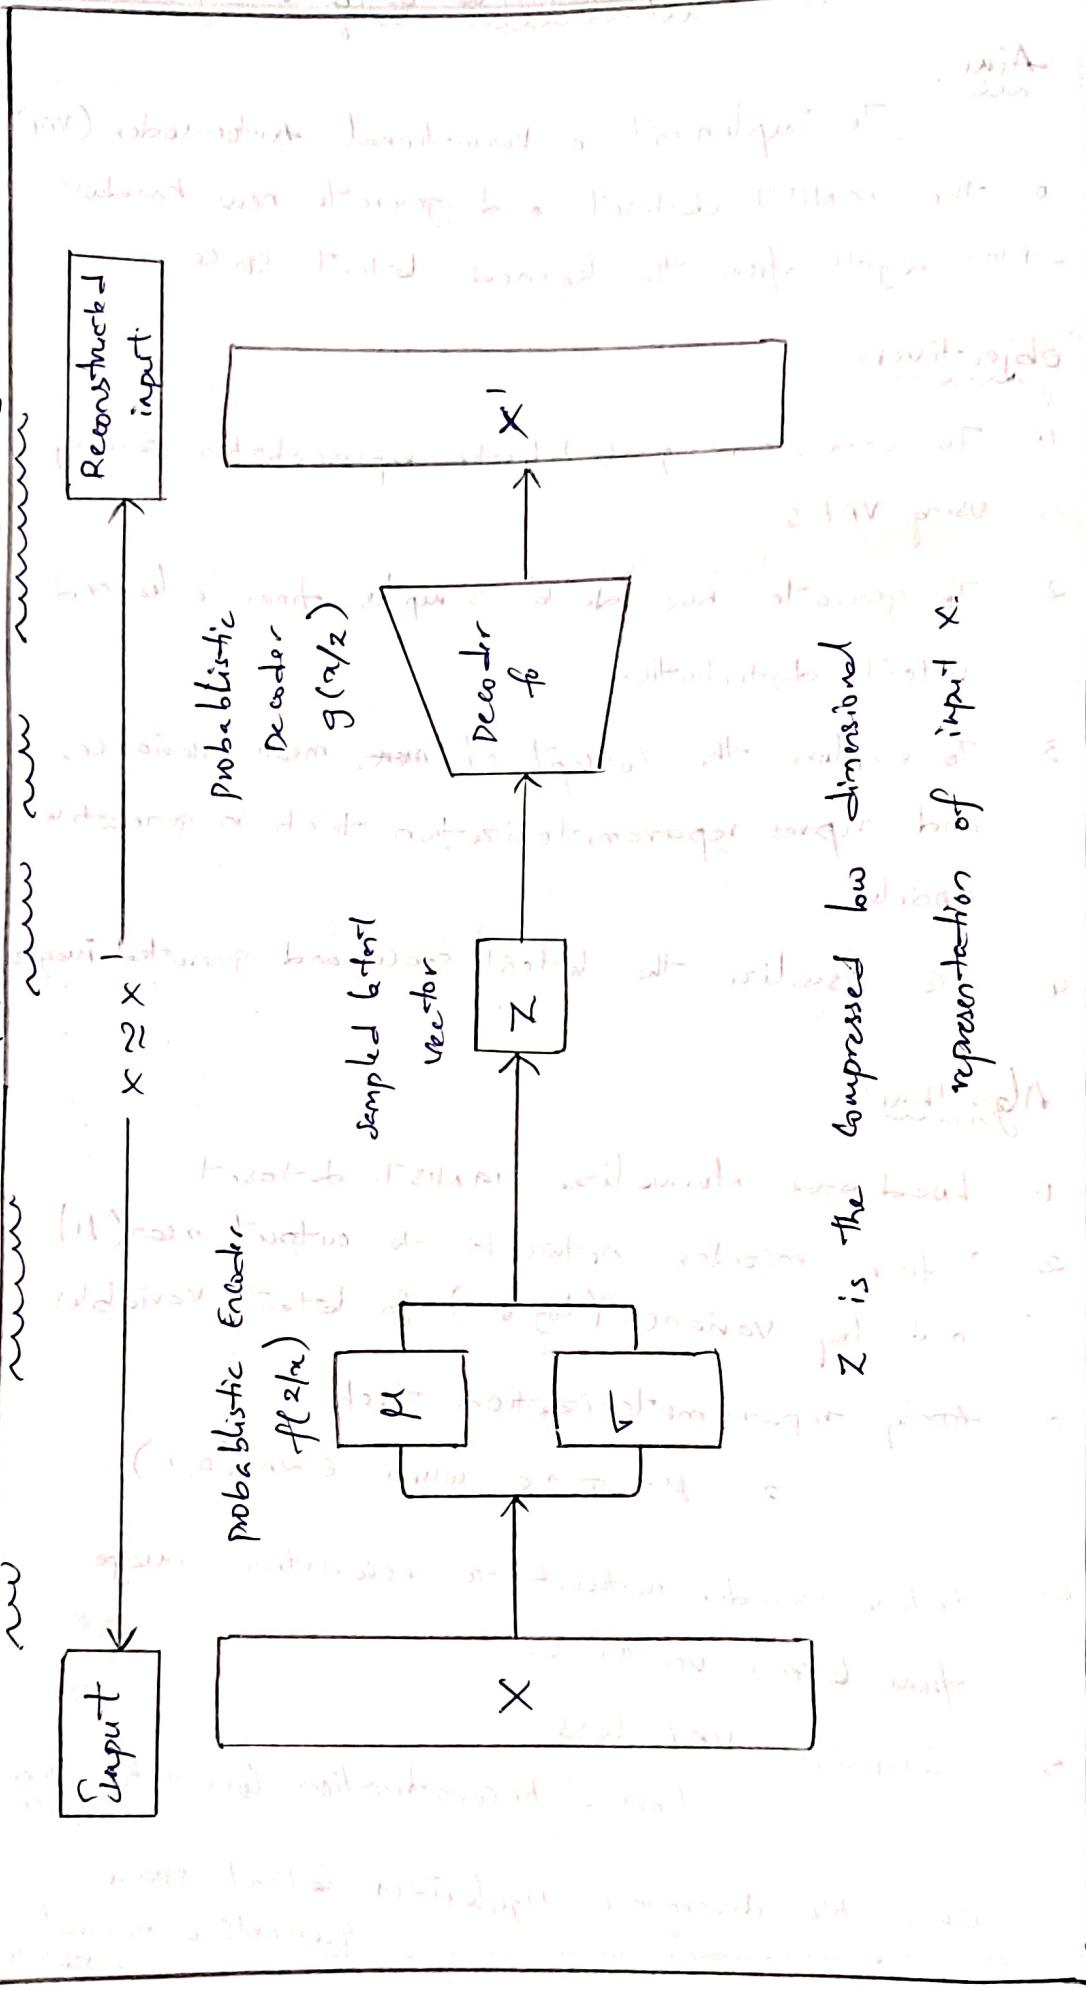
Observations

- The VAE learns a smooth, continuous latent space where similar digits are close to each other.
- Random sampling in latent space generates realistic digits.

Conclusion:

VAEs enhance autoencoders by adding probabilistic latent space, enabling both data compression & generation.

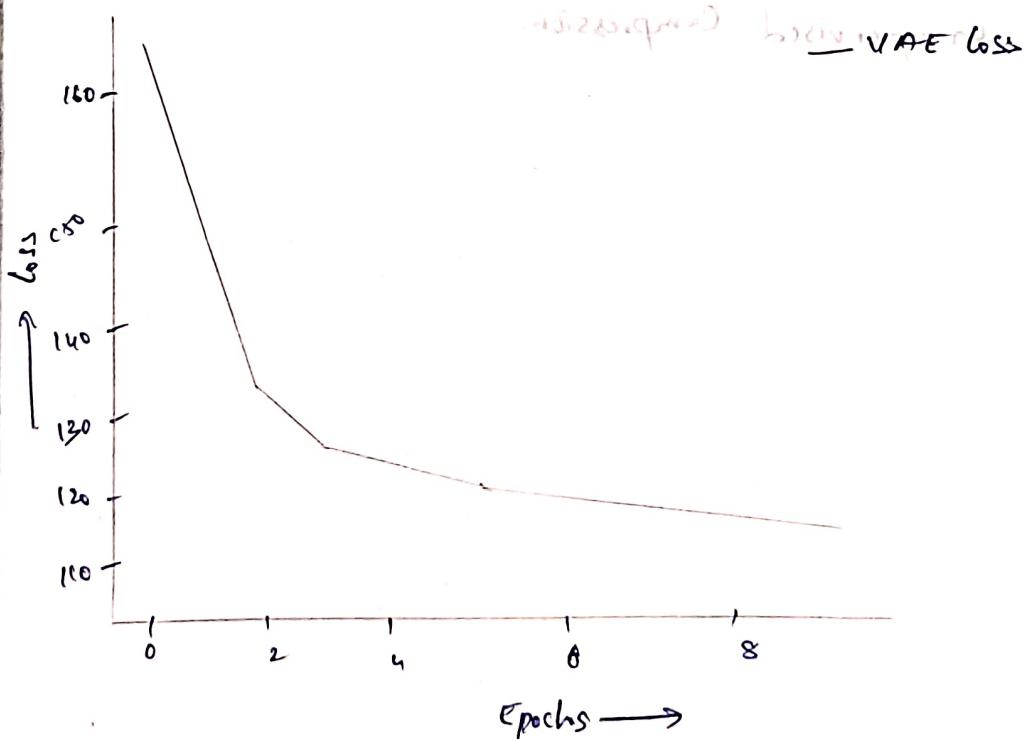
VAE BASIC ARCHITECTURE.



Output!

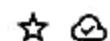
Epoch [1/10] Loss: 164.3748
Epoch [2/10] Loss: 121.2820
Epoch [3/10] Loss: 114.5211
Epoch [4/10] Loss: 111.6053
Epoch [5/10] Loss: 109.8342
Epoch [6/10] Loss: 108.6767
Epoch [7/10] Loss: 107.8205
Epoch [8/10] Loss: 107.1637
Epoch [9/10] Loss: 106.6910
Epoch [10/10] Loss: 106.2222

Training loss of VAE





dltlab11.ipynb



File Edit View Insert Runtime Tools Help

Commands + Code + Text | ▶ Run all ▾



```
# vae_mnist.py
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# ---- Device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# ---- Load MNIST dataset
tf = transforms.Compose([transforms.ToTensor()])
train_data = datasets.MNIST(root='data', train=True, transform=tf, download=True)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)

# ---- Variational Autoencoder
class VAE(nn.Module):
    def __init__(self):
        super(VAE, self).__init__()
        self.fc1 = nn.Linear(784, 400)
        self.fc_mu = nn.Linear(400, 20)      # Mean of latent vector
        self.fc_logvar = nn.Linear(400, 20) # Log variance of latent vector
        self.fc2 = nn.Linear(20, 400)
        self.fc3 = nn.Linear(400, 784)
        self.relu = nn.ReLU()
        self.sigmoid = nn.Sigmoid()

    def encode(self, x):
        h1 = self.relu(self.fc1(x))
        mu = self.fc_mu(h1)
        logvar = self.fc_logvar(h1)
        return mu, logvar

    def reparameterize(self, mu, logvar):
        std = torch.exp(0.5 * logvar)
        eps = torch.randn_like(std)
        return mu + eps * std # z = μ + σ * ε

    def decode(self, z):
        h2 = self.relu(self.fc2(z))
        return self.sigmoid(self.fc3(h2))
```

```
def forward(self, x):
    x = x.view(-1, 784)
    mu, logvar = self.encode(x)
    z = self.reparameterize(mu, logvar)
    recon = self.decode(z)
    return recon, mu, logvar

# ---- Loss Function (Reconstruction + KL Divergence)
def vae_loss(recon_x, x, mu, logvar):
    BCE = nn.functional.binary_cross_entropy(recon_x, x.view(-1, 784), reduction='sum')
    # KL Divergence: how latent distribution differs from standard normal
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

# ---- Model, Optimizer
model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=1e-3)

# ---- Training
epochs = 10
train_losses = []

for epoch in range(1, epochs + 1):
    model.train()
    total_loss = 0
    for imgs, _ in train_loader:
        imgs = imgs.to(device)
        optimizer.zero_grad()
        recon, mu, logvar = model(imgs)
        loss = vae_loss(recon, imgs, mu, logvar)
        loss.backward()
        total_loss += loss.item()
        optimizer.step()

    avg_loss = total_loss / len(train_loader.dataset)
    train_losses.append(avg_loss)
    print(f"Epoch [{epoch}/{epochs}] Loss: {avg_loss:.4f}")

# ---- Plot Training Loss
plt.plot(train_losses, label="VAE Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
```

```
# ---- Plot Training Loss
plt.plot(train_losses, label="VAE Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss of Variational Autoencoder")
plt.legend()
plt.show()

# ---- Reconstruction Visualization
model.eval()
imgs, _ = next(iter(train_loader))
imgs = imgs[:8].to(device)
with torch.no_grad():
    recons, _, _ = model(imgs)

imgs, recons = imgs.cpu(), recons.view(-1, 1, 28, 28).cpu()

fig, axes = plt.subplots(2, 8, figsize=(12, 3))
for i in range(8):
    axes[0, i].imshow(imgs[i].squeeze(), cmap='gray')
    axes[0, i].axis('off')
    axes[1, i].imshow(recons[i].squeeze(), cmap='gray')
    axes[1, i].axis('off')
plt.suptitle("Top: Original Images | Bottom: VAE Reconstructions")
plt.show()

# ---- Sample new images from random latent space
with torch.no_grad():
    z = torch.randn(16, 20).to(device)
    samples = model.decode(z).view(-1, 1, 28, 28).cpu()
    fig, axes = plt.subplots(2, 8, figsize=(12, 3))
    for i in range(16):
        axes[i//8, i%8].imshow(samples[i].squeeze(), cmap='gray')
        axes[i//8, i%8].axis('off')
    plt.suptitle("Generated Digits from Random Latent Space")
    plt.show()
```

100% | [██████████] | 9.91M/9.91M [00:02<00:00, 4.93MB/s]
100% | [██████████] | 28.9k/28.9k [00:00<00:00, 131kB/s]
100% | [██████████] | 1.65M/1.65M [00:01<00:00, 1.24MB/s]
100% | [██████████] | 4.54k/4.54k [00:00<00:00, 10.3MB/s]

Epoch [1/10] Loss: 164.3748
Epoch [2/10] Loss: 121.2820
Epoch [3/10] Loss: 114.5211
Epoch [4/10] Loss: 111.6053
Epoch [5/10] Loss: 109.8342
Epoch [6/10] Loss: 108.6767
Epoch [7/10] Loss: 107.8205
Epoch [8/10] Loss: 107.1637
Epoch [9/10] Loss: 106.6910
Epoch [10/10] Loss: 106.2222

Training Loss of Variational Autoencoder

