

GAN to generate complex color images

Exp No: 12 :- Implement a Deep Convolutional

Aim:- To implement a Deep Convolutional Generative Adversarial Network Capable of generating realistic color images.

Objectives:-

1. To understand the working of GANs with generator and discriminator networks
2. To build a DCGAN using convolutional and transpose convolutional layers
3. To train the model to generate color images from random noise
4. To visualize and evaluate generated image quality

Algorithm:-

1. Import and preprocess color image dataset. (CIFAR-10)
2. Normalize image data to range $[-1, 1]$.
3. Build the Generator.
→ Use Conv2DTranspose layers with Batch Norm & ReLU.
→ output 3-channel (RGB) image using tanh activation
4. Build the Discriminator.
→ Use Conv2D layers with Leaky ReLU & Dropout.
→ output binary classification (real/fake)
5. Train both networks adversarially:
→ Discriminator learns to distinguish real/fake
→ Generator learns to fool the discriminator
6. Generate & visualize synthetic images.

Pseudo Code:

Load CIFAR-10 dataset

normalize images to $[-1, 1]$

Define Generator:

Dense \rightarrow Reshape \rightarrow Conv2D Transpose (ReLU)

output: Conv2D Transpose (3, tanh).

Define Discriminator:

Conv2D \rightarrow Leaky ReLU \rightarrow Dropout

output: Dense(1, sigmoid)

Combine models into DCGAN

Train:

For each epoch:

1. Train discriminator with Real + fake image
2. Train generator via combined model

Generate and visualize images.

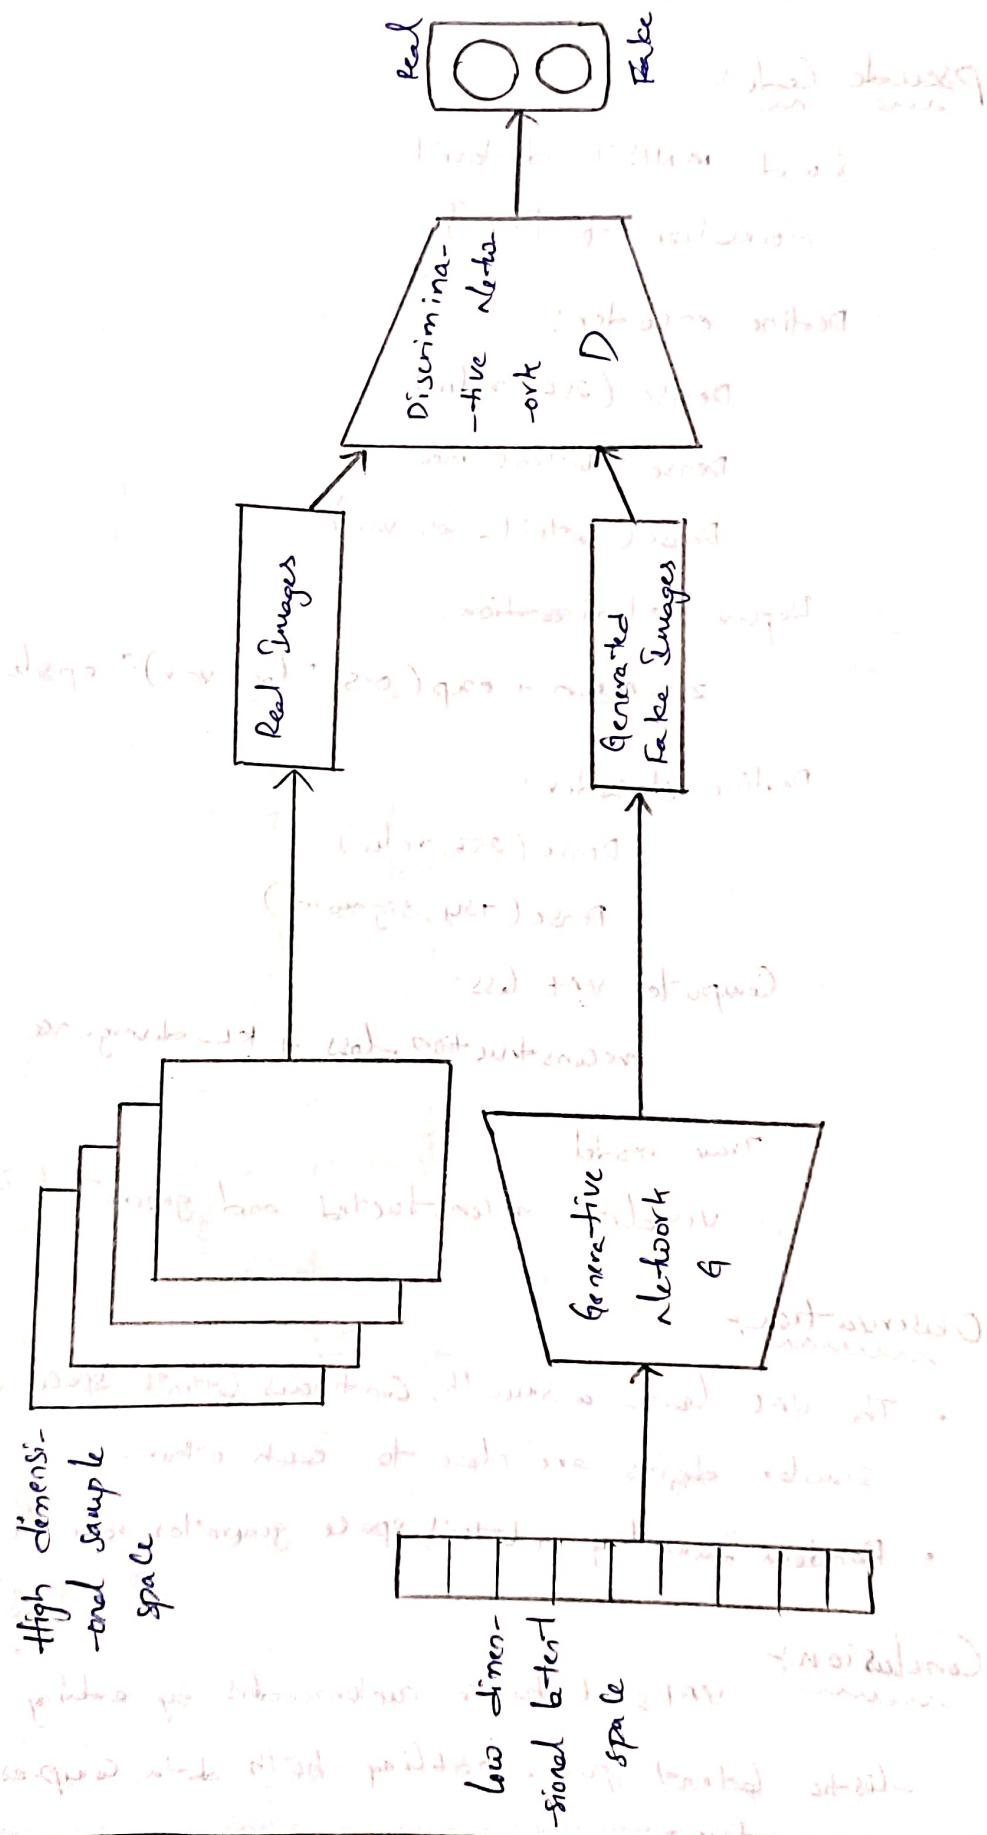
Observation:

- The DCGAN learns to produce increasingly realistic color images as training progresses.
- Generated images evolve from random noise to clear structures.

Conclusion:

DCGANs effectively learn to generate complex, realistic color images using adversarial training b/w the generator & discriminator. They demonstrate the power of deep convolutional networks for unsupervised image generation.

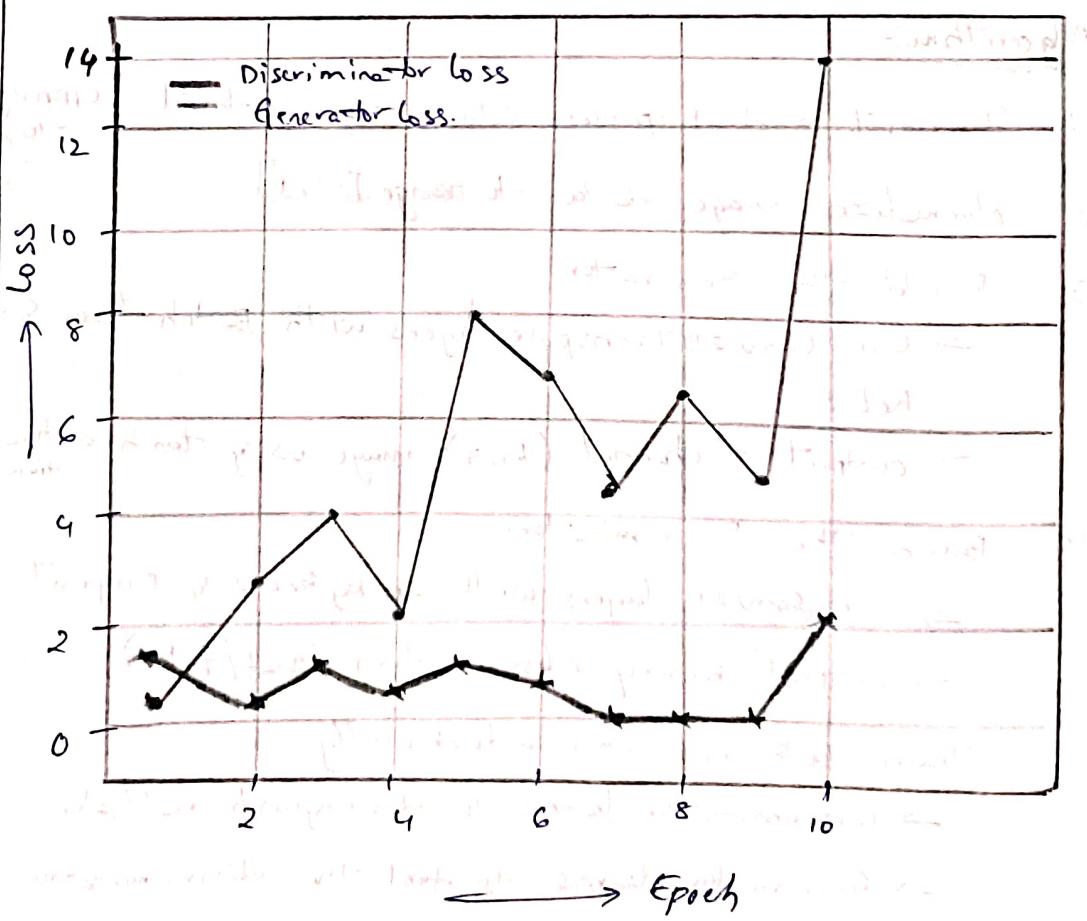
GAN ARCHITECTURE FOR IMAGE



Output:-

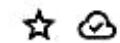
Epoch [0/10] D loss: 1.0992, G loss: 0.5678
Epoch [1/10] D loss: 0.5060, G loss: 2.7113
Epoch [2/10] D loss: 1.2431, G loss: 3.9227
Epoch [3/10] D loss: 0.4229, G loss: 2.2684
Epoch [4/10] D loss: 1.1486, G loss: 3.1007
Epoch [5/10] D loss: 1.0081, G loss: 6.7625
Epoch [6/10] D loss: 0.0252, G loss: 4.5714
Epoch [7/10] D loss: 0.0043, G loss: 6.3371
Epoch [8/10] D loss: 0.1380, G loss: 4.7357
Epoch [9/10] D loss: 0.3294, G loss: 14.0363

Epoch vs Loss (GAN).





dltlab12.ipynb



File Edit View Insert Runtime Tools Help

Commands + Code + Text | ▶ Run all ▾



[]

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import torchvision.utils as vutils
import matplotlib.pyplot as plt
import numpy as np

# 1. Hyperparameters
batch_size = 128
latent_dim = 100
epochs = 10
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# 2. Data loading (CIFAR-10 color images)
transform = transforms.Compose([
    transforms.Resize(64),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
dataset = datasets.CIFAR10(root='./data', download=True, transform=transform)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)

# 3. Define Generator
class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()
        self.model = nn.Sequential(
            nn.ConvTranspose2d(100, 512, 4, 1, 0, bias=False),
            nn.BatchNorm2d(512), nn.ReLU(True),
            nn.ConvTranspose2d(512, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256), nn.ReLU(True),
            nn.ConvTranspose2d(256, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128), nn.ReLU(True),
            nn.ConvTranspose2d(128, 64, 4, 2, 1, bias=False),
            nn.BatchNorm2d(64), nn.ReLU(True),
            nn.ConvTranspose2d(64, 3, 4, 2, 1, bias=False),
            nn.Tanh()
        )
    def forward(self, x):
        return self.model(x)
```



```
# 4. Define Discriminator
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.model = nn.Sequential(
            nn.Conv2d(3, 64, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(64, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128), nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(128, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256), nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(256, 512, 4, 2, 1, bias=False),
            nn.BatchNorm2d(512), nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(512, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )
    def forward(self, x):
        return self.model(x).view(-1, 1)

# 5. Initialize models
netG, netD = Generator().to(device), Discriminator().to(device)
criterion = nn.BCELoss()
optimizerG = optim.Adam(netG.parameters(), lr=0.0002, betas=(0.5, 0.999))
optimizerD = optim.Adam(netD.parameters(), lr=0.0002, betas=(0.5, 0.999))

# Lists to store losses
d_losses = []
g_losses = []

# 6. Training loop
for epoch in range(epochs):
    for i, (real_imgs, _) in enumerate(dataloader):
        real_imgs = real_imgs.to(device)
        batch_size = real_imgs.size(0)

        # Real labels = 1, Fake labels = 0
        real_labels = torch.ones(batch_size, 1, device=device)
        fake_labels = torch.zeros(batch_size, 1, device=device)

        # Train Discriminator
        z = torch.randn(batch_size, latent_dim, 1, 1, device=device)
        fake_imgs = netG(z)
```

```
real_loss = criterion(netD(real_imgs), real_labels)
fake_loss = criterion(netD(fake_imgs.detach()), fake_labels)
d_loss = real_loss + fake_loss

optimizerD.zero_grad()
d_loss.backward()
optimizerD.step()

# Train Generator
g_loss = criterion(netD(fake_imgs), real_labels)
optimizerG.zero_grad()
g_loss.backward()
optimizerG.step()

# Store losses
d_losses.append(d_loss.item())
g_losses.append(g_loss.item())

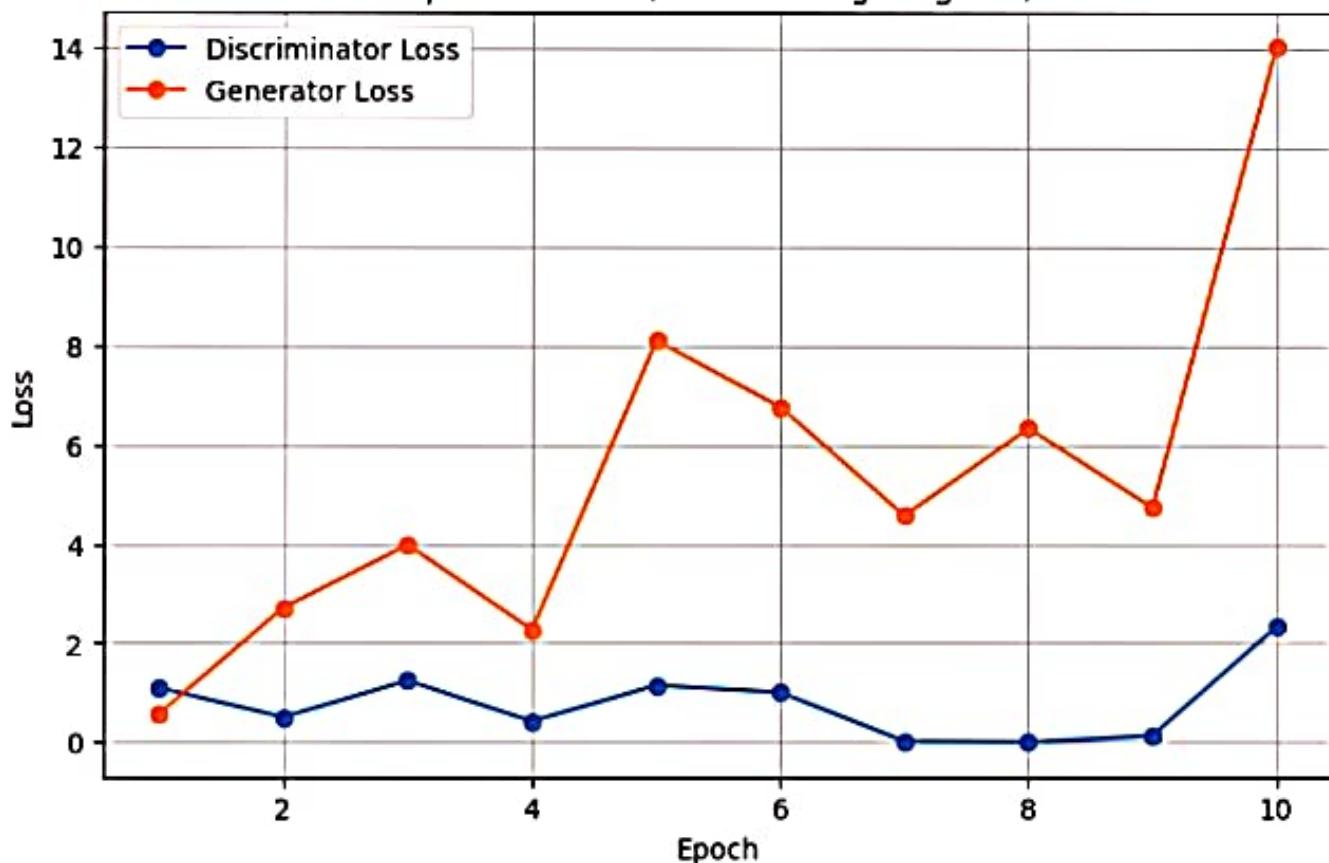
print(f"Epoch [{epoch+1}/{epochs}] D Loss: {d_loss:.4f}, G Loss: {g_loss:.4f}")

# 7. Plot Epoch Graph (Loss vs Epoch)
plt.figure(figsize=(8,5))
plt.plot(range(1, epochs+1), d_losses, label='Discriminator Loss', marker='o')
plt.plot(range(1, epochs+1), g_losses, label='Generator Loss', marker='o')
plt.title("Epoch vs Loss (GAN Training Progress)")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()

# 8. Generate samples
z = torch.randn(64, latent_dim, 1, 1, device=device)
fake_imgs = netG(z)
vutils.save_image(fake_imgs, 'dcgan_generated.png', normalize=True)
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Generated Images")
plt.imshow(np.transpose(vutils.make_grid(fake_imgs[:64], padding=2, normalize=True).cpu(), (1,2,0)))
plt.show()
```

100%|██████████| 170M/170M [00:14<00:00, 11.5MB/s]
Epoch [1/10] D Loss: 1.0992, G Loss: 0.5678
Epoch [2/10] D Loss: 0.5060, G Loss: 2.7113
Epoch [3/10] D Loss: 1.2431, G Loss: 3.9827
Epoch [4/10] D Loss: 0.4229, G Loss: 2.2684
Epoch [5/10] D Loss: 1.1486, G Loss: 8.1007
Epoch [6/10] D Loss: 1.0081, G Loss: 6.7625
Epoch [7/10] D Loss: 0.0252, G Loss: 4.5718
Epoch [8/10] D Loss: 0.0043, G Loss: 6.3371
Epoch [9/10] D Loss: 0.1380, G Loss: 4.7357
Epoch [10/10] D Loss: 2.3294, G Loss: 14.0363

Epoch vs Loss (GAN Training Progress)



Generated Images

